

Rainbow: Engineering Support for Self-Healing Systems

CompArch 2009

David Garlan
June 2009



Joint work with ...



- Owen Cheng
 - PhD Thesis May 2008, now at NASA JPL
- Bradley Schmerl
 - Staff
- Jung Soo Kim
 - Current PhD student

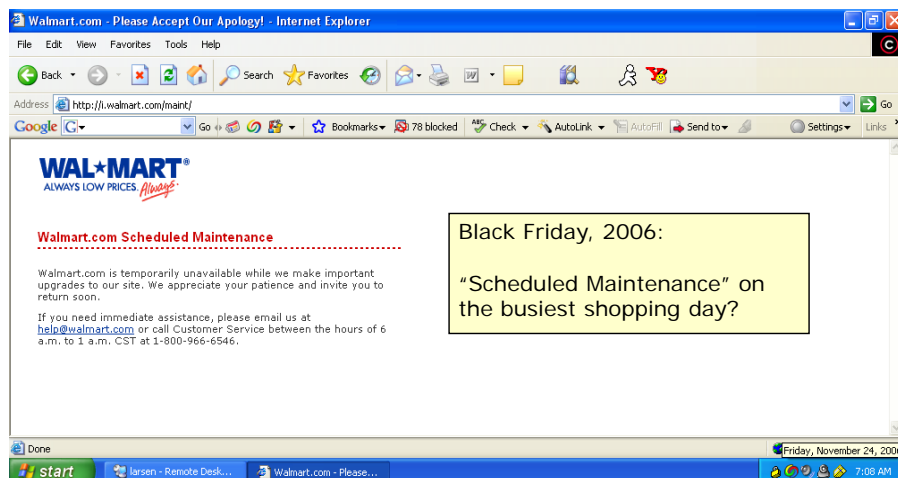
The Problem

- An important requirement for modern software-based systems

Maintain high-availability and optimal performance even in the presence of

- changes in environment
- system faults
- changes in user needs and context

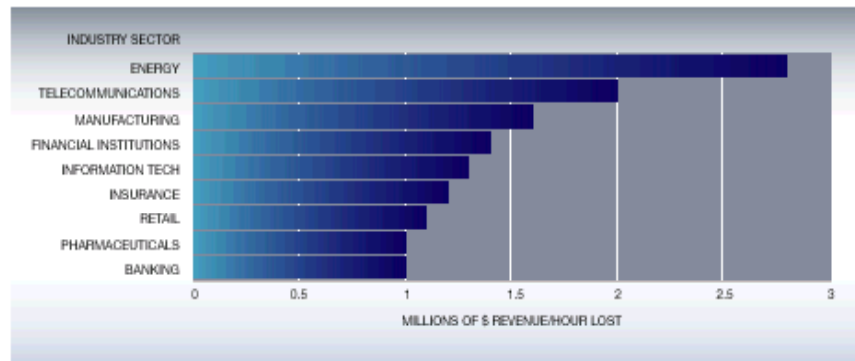
Websites Fail to Adapt



- Amazon.com disrupted due to Xbox 360 on the day before

Cost of Downtime

- Average hourly impact of downtime by industry sector



Data from *IT Performance Engineering and Measurement Strategies: Quantifying Performance Loss*, Meta Group, Stamford, CT (October 2000).

© David Garlan 2009

5

How is this addressed today?

- **Technique 1:** try to build resilience into application code directly
 - Use exceptions, timeouts, and other low-level programming mechanisms
- Unfortunately, this approach is not good for
 - Handling changing objectives
 - Locating the true cause of the problem
 - Detecting “softer” system anomalies
 - Maintainability: hard to add and modify adaptation policies and mechanisms
 - Legacy systems: hard to retrofit later

© David Garlan 2009

6

How is this addressed today?

- **Technique 2:** human oversight
 - Operators, system administrators
 - Global oversight, intelligent response
- Unfortunately, this approach is
 - Costly
 - Error-prone

Cost of Human Oversight


- Estimated 1/3-1/2 total IT budget to prevent or recover from crash
- “For every dollar to purchase storage, you spend \$9 to have someone manage it”—Nick Tabellion
- Administrative cost: 60-75% overall cost of database ownership
- 40% of root causes of computer system outage is attributable to operator error

A New Approach

- Goal: systems automatically and optimally adapt to handle
 - changes in user needs
 - variable resources
 - faults

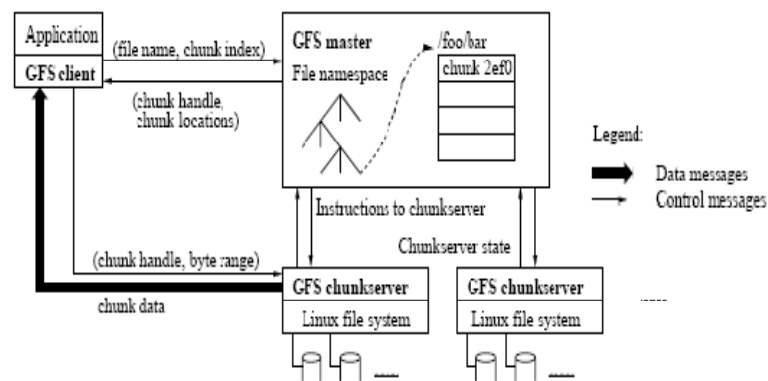
But how?

Answer: Move from open-loop to closed-loop systems



© David Garlan 2009

Example: Google File System



Source: "The Google File System" Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. SOSP 2003.

Figure 1: GFS Architecture

© David Garlan 2009

10

The Challenge

- Engineer self-adaptation to support
 - Cost-effectiveness
 - Legacy systems
 - Domain-specific adaptations
 - Multiple quality dimensions
 - Ease of changing adaptation policies


This Talk

- *Software Architectures* for dynamic self-adaptation
 - *Rainbow*: a framework in which architectural models can be used to adapt systems
 - *Stitch*: a new language to express self-adaptation strategies
- Examples and evaluation
- On-going work
- Related work

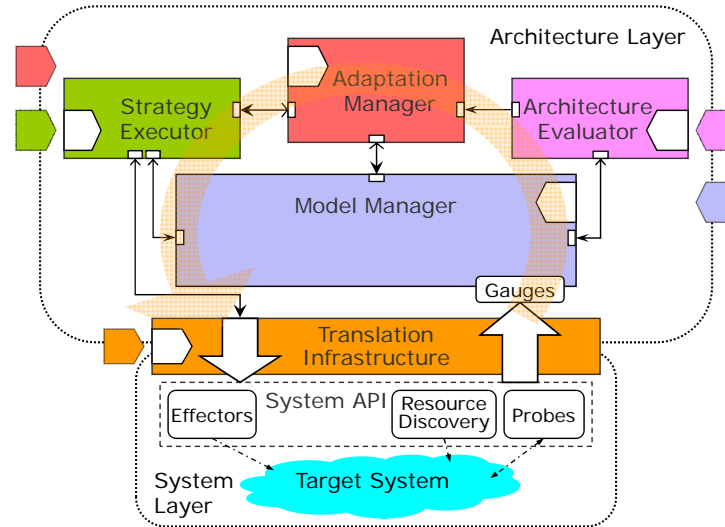
Our Approach

- **Framework**
 - Called Rainbow
 - Common infrastructure with reusable mechanisms
 - Customization points for tailoring to the domain
- **Language**
 - Called Stitch
 - Expresses adaptation policies as *strategies*
 - Expresses business quality dimensions and preferences
- **Adaptation engineering process**
 - Process for customizing the Rainbow framework
 - Roles: for engineering self-adaptation capabilities

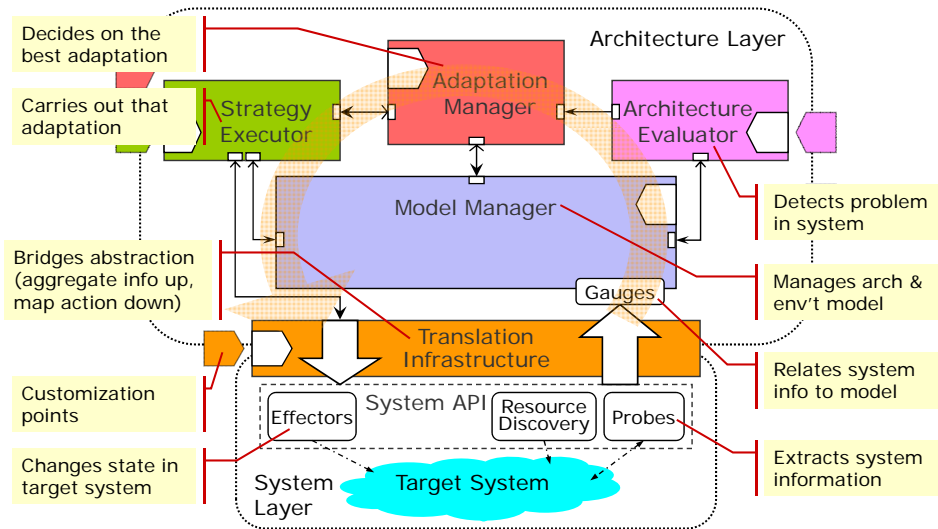
The Rainbow Framework



Rainbow Framework Overview




Rainbow Framework Overview



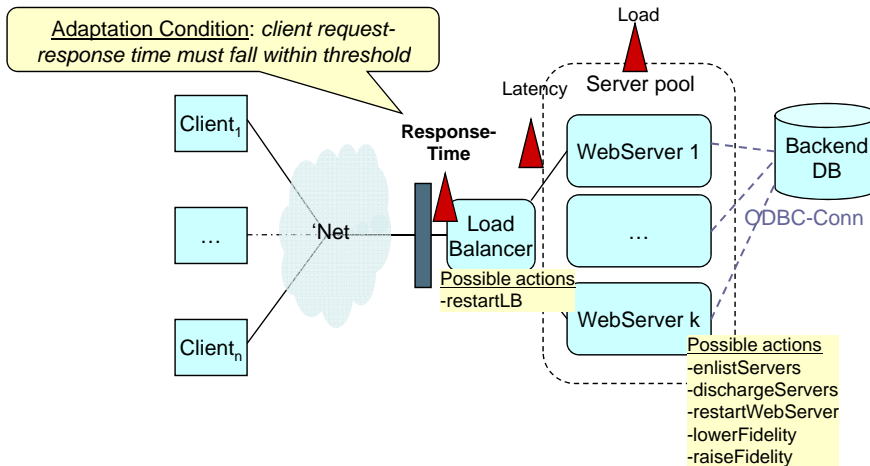
Self-Adaptation Example: *Znn.com*

- We'll use this example throughout our talk

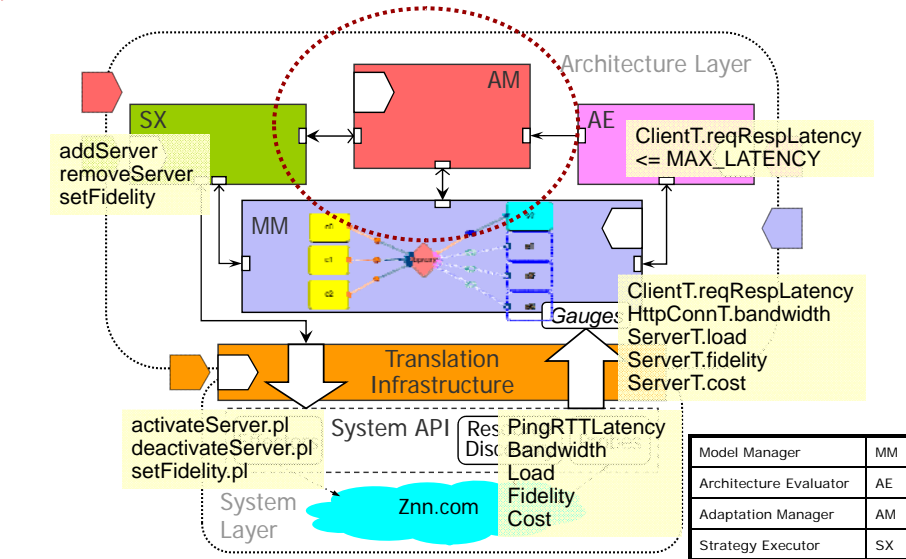


Self-Adaptation Example: *Znn.com*

- We'll use this example throughout our talk



Znn.com: Rainbow Customizations (Part 1)

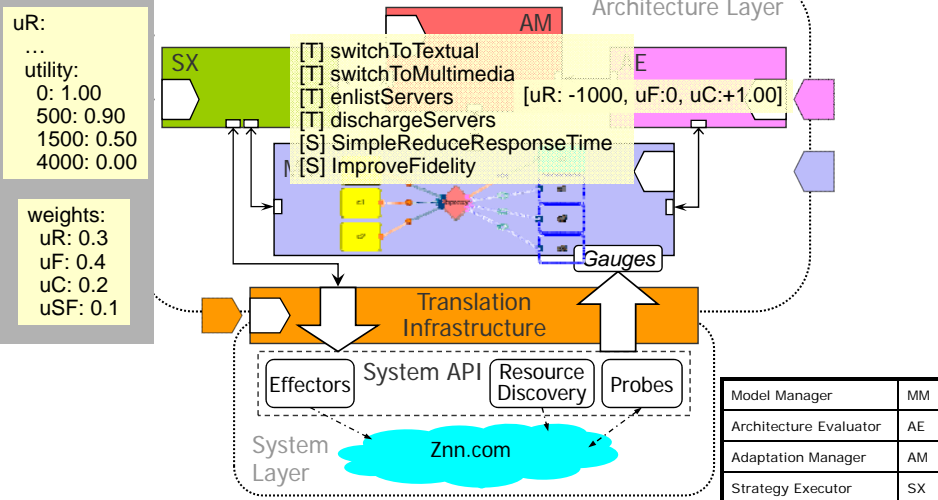


© David Garlan 2009

19

Znn.com: Rainbow Customizations (Part 2)

Objectives: timely response (uR), high-quality content (uF), low-provision cost (uC)



© David Garlan 2009

20

Stitch Language Overview

- Language requirements
 - **Concepts**: concepts to formalize operational aspects of self-adaptation
 - **Value system**: way to specify value system for comparing adaptations
 - **Choice**: apply value system to select best course of adaptation
- **Concepts** to express adaptation knowledge
 - *Operator* – basic system-provided command
 - *Tactic* – action with impact on quality dimensions
 - *Strategy* – adaptation with intermediate steps of condition-action-delay
- A way to specify **value system**
 - *Quality dimensions* – business QoS concerns
 - *Utility preferences* – business priorities over dimensions
 - *Adaptation conditions* – opportunities for improvement
- **Choice** of best adaptation
 - Strategy selection – best adaptation for current system conditions

Innovative features of Stitch

- Control-system model: Selection of next action in a strategy depends on observed effects of previous action
- Value system: Utility-based selection of best strategy allows context-sensitive adaptation
- Asynchrony: Timing delays capture settling time
- Uncertainty: effectiveness of a given tactic is known only within some probability

Stitch Language Constructs

Constructs

- **Operator**
- Tactic
- Strategy
- Quality dimensions
- Utility preferences
- Adaptation conditions

- A basic command provided by target system
 - Defined in the architectural style
 - Mapped to system-level effectors

```

addServer
removeServer
setFidelity
↓
activateServer.pl
deactivateServer.pl
setFidelity.pl
    
```

Stitch Language Constructs

Constructs

- Operator
- **Tactic**
- Strategy
- Quality dimensions
- Utility preferences
- Adaptation conditions

- Provides an **abstraction** to
 - **Package** operators into larger units
 - Form primitive adaptation **step**
 - Associate **cost-benefit** impact on quality dimensions
- Characteristics
 - An action composed of a sequence of operators
 - Guard conditions to determine applicability
 - Expected effects to observe in target system
 - Specified impacts on quality dimensions
 - Execution context: architectural style (types + operators) and model instance
 - Read-only snapshot of model during execution
 - Uses operators, provides primitives for *strategies*
- Cost-benefit attributes
 - $[a_1:\Delta v_1, \dots, a_n:\Delta v_n]$ $[uR: -1000, uF: 0, uC: +1.00]$

Stitch Language Constructs

Constructs

- Operator
- **Tactic** (example)
- Strategy

```
module newssite.tactics.example;
import model "ZnnSys.acme" { ZnnSys as M, ZnnFam as T };
import op "newssite.operator.ArchOp" { ArchOp as Sys };

tactic switchToTextualMode () {
  condition {
    exists c : T.ClientT in M.components | c.experRespTime >
      M.MAX_RESPTIME;
  }
  action {
    svrs = { select s : T.ServerT | !s.isTextualMode };
    for (T.ServerT s : svrs) {
      Sys.setTextualMode(s, true);
    }
  }
  effect {
    forall c : T.ClientT in M.components | c.experRespTime <=
      M.MAX_RESPTIME;
    forall s : T.ServerT in M.components | s.isTextualMode;
  }
}
```

Stitch Language Constructs

Constructs

- Operator
- **Tactic** (example)
- Strategy

Import of model
and operators

Conditions of
applicability

Sequence of
operators

Expected effects
to observe

```
module newssite.tactics.example;
import model "ZnnSys.acme" { ZnnSys as M, ZnnFam as T };
import op "newssite.operator.ArchOp" { ArchOp as Sys };

tactic switchToTextualMode () {
  condition {
    exists c : T.ClientT in M.components | c.experRespTime >
      M.MAX_RESPTIME;
  }
  action {
    svrs = { select s : T.ServerT | !s.isTextualMode };
    for (T.ServerT s : svrs) {
      Sys.setTextualMode(s, true);
    }
  }
  effect {
    forall c : T.ClientT in M.components | c.experRespTime <=
      M.MAX_RESPTIME;
    forall s : T.ServerT in M.components | s.isTextualMode;
  }
}
```

Stitch Language Constructs

Constructs

- Operator
- Tactic
- **Strategy**
- Quality dimensions
- Utility preferences
- Adaptation conditions

- Encapsulates path of adaptations where
 - Each step is conditional execution of a tactic
 - Choice of tactic depend on intermediate results
 - Intermediate step might reduce utility
 - Timing delay for asynchrony in achieving effect
 - Probabilities on conditions to handle inherent uncertainties in achieving effect
- Characteristics
 - Tree of condition-action-delay decision nodes
 - Conditions of applicability to determine involvement
 - Execution context: arch style (types) and tactics
 - Intermediate system observations at decision nodes
 - Enables computing aggregate cost-benefit attributes for utility-based strategy selection

Stitch Language Constructs

Constructs

- Operator
- Tactic
- **Strategy** (e.g.)

```

module newssite.strategies.example;
import model "ZnnSys.acme" { ZnnSys as M, ZnnFam as T };
import lib "newssite.tactics.example";
import op "org.sa.rainbow.stitch.lib.*"; // Model, Set, & Util

define boolean styleApplies = ... Model.hasType(M, "ServerT");
define boolean cViolation = exists c : T.ClientT in M.components |
    c.experRespTime > M.MAX_RESPTIME;

strategy SimpleReduceResponseTime [ styleApplies && cViolation ] {
    define boolean hiLatency = exists conn : T.HttpConnT in
        M.connectors | conn.latency > M.MAX_LATENCY;
    define boolean hiLoad = exists s : T.ServerT in M.components |
        s.load > M.MAX_UTIL;

    t1: (#[Pr{t1}] hiLatency) -> switchToTextualMode() @[1000/*ms*/] {
        t1a: (success) -> done ;
    }
    t2: (#[Pr{t2}] hiLoad) -> enlistServer(1) @[2000/*ms*/] {
        t2a: (!hiLoad) -> done ;
        t2b: (!success) -> do [1] t1 ;
    }
    t3: (default) -> fail;
}
    
```

Stitch Language Constructs

Constructs

- Operator
- Tactic
- Strategy (e.g.)

Import of model and tactics

Condition of applicability

Step of condition-action-delay

Timing delay

Tactic succeeded?

Complete strategy

No branch matches?

Abort strategy

```
module newssite.strategies.example;
import model "ZnnSys.acme" { ZnnSys as M, ZnnFam as T };
import/lib "newssite.tactics.example";
import op "org.sa.rainbow.stitch.lib.*"; // Model, Set, & Util

define boolean styleApplies = ... Model.hasType(M, "ServerT");
define boolean cViolation = exists c : T.ClientT in M.components |
c.experRespTime > M.MAX_RESPTIME;

strategy SimpleReduceResponseTime [ styleApplies && cViolation ] {
  define boolean hiLatency = exists conn : T.HttpConnT in
M.connectors | conn.latency > M.MAX_LATENCY;
  define boolean hiLoad = exists s : T.ServerT in M.components |
s.load > M.MAX_UTIL;

  t1: (#[Pr{t1}] hiLatency) -> switchToTextualMode() @[1000/*ms*/] {
    t1a: (success) -> done ;
  }
  t2: (#[Pr{t2}] hiLoad) -> enlistServer(1) @[2000/*ms*/] {
    t2a: (!hiLoad) -> done ;
    t2b: (!success) -> do [1] t1 ;
  }
  t3: (default) -> fail;
}
```

Stitch Language Constructs

Constructs

- Operator
- Tactic
- Strategy
- Quality dimensions
- Utility preferences
- Adaptation conditions

- Intuitively, determine *what* to adapt for
- Notion of utility for particular values of attribute
- Profile: ID, label, description, mapping, function
 - Mapping: monitored architectural property → current system condition
 - Function: linear, sigmoid, custom (interpolated)

```
uR:
label: "Average Response Time"
mapping: "[EAvg] ClientT.experRespTime"
description: ...
utility:
0: 1.00
500: 0.90
1500: 0.50
4000: 0.00
```

Stitch Language Constructs

Constructs

- Operator
- Tactic
- Strategy
- Quality dimensions
- **Utility preferences**
- Adaptation conditions

- Intuitively, determine *how* to adapt
- Resource constraints: can't optimize all qualities
- Technique: assign weights for relative importance
 - Weights must sum to 1

```
weights:  
uR: 0.3  
uF: 0.4  
uC: 0.2  
uSF: 0.1
```

Stitch Language Constructs

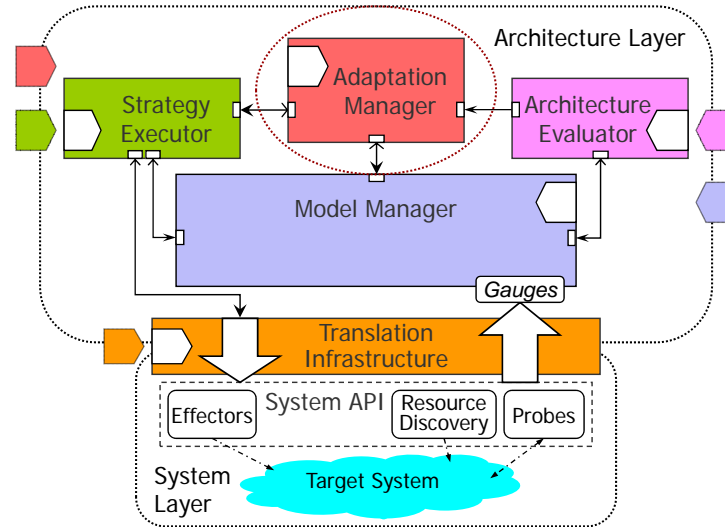
Constructs

- Operator
- Tactic
- Strategy
- Quality dimensions
- Utility preferences
- **Adaptation conditions**

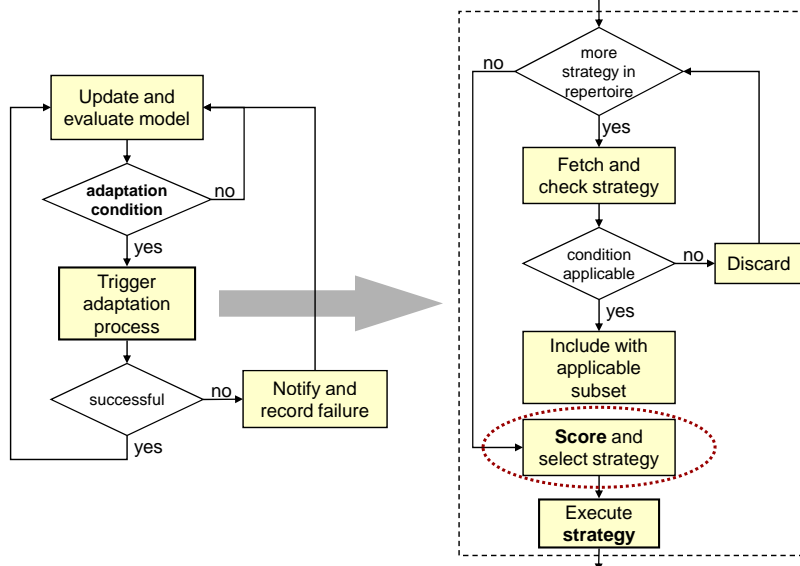
- Intuitively, determine *when* to adapt
- Constraints defined in architectural style
- QoS threshold quantity defined in model instance

```
Invariant self.avg_latency < MAX_RESPTIME;  
Property MAX_RESPTIME : float = 1000.0;
```


Rainbow Context for Adaptation Execution



Adaptation Manager: Adaptation Execution



Strategy Selection

- Given:
 - Quality dimensions and weights (e.g., 4)
 - A strategy with
 - N nodes
 - Branch probabilities as shown
 - Tactic cost-benefit attributes

$$\begin{aligned}
 &u_{latency}(), u_{quality}(), u_{cost}(), u_{disruption}() \\
 &(w_{latency}, w_{quality}, w_{cost}, w_{disruption}) \\
 &= (0.5, 0.3, 0.1, 0.1) [= 1]
 \end{aligned}$$

Algorithm


Given tree g with node x and its children c :

$$EAV(g) = sysAV + AggAV(\text{root}(g))$$

$$AggAV(x) = cbav(x) + \sum_c prob(x,c) AggAV(c)$$

- Propagate cost-benefit vectors up the tree, reduced by branch probabilities
- Merge expected vector with current conditions (assume: [1125, 3.5, 0, 0])
- Evaluate quality attributes against utility functions
- Compute weighted sum to get utility **score**

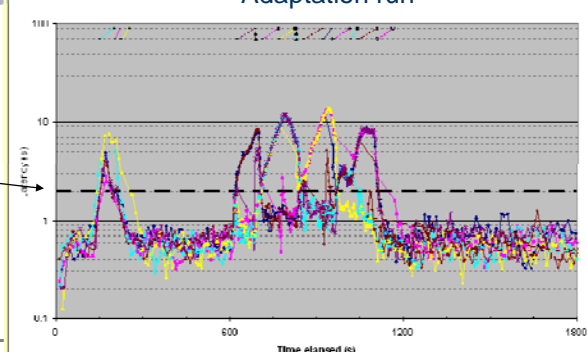
Control run



System Adapts

Data shows that our adaptation approach improves overall system performance


Adaptation run



Latency = 2 secs

Znn.com Experiment Run Results

Aggregate JMeter Data									
URL	#Samples	Average	Median	90% Line	Min	Max	Error%	Throughput	KB/sec
0415K No Rainbow	1200	7981	1953	37514	93	52201	0.00%	8.4/sec	6.92
0415L SimpleReduce1	1200	1502	16	2187	0	22202	0.17%	29.5/sec	78.26




© David Garlan 2009

37

System Administration Examples

- Sys-admin interview
 - **Results:** Stitch concepts seem **natural** fit for sys-admin routines
 - Methodology: priming, interview, compose Stitch script from scenarios
 - White problem scenarios: scripts represented in Stitch
 - Almossawi problem scenarios: structure matches Stitch strategies
- Analysis using CMU sys-admin example: *Netbwe*
 - **Results:**
 - Rainbow captured adaptation concerns
 - Stitch **hoisted** policies buried in Perl code
 - Distinguishable adaptation tasks
 - Core commands as operators
 - Coarser-grained sequence of commands (step) with conditions of applicability and intended effects
 - Adaptations with intermediate condition-actions and observations



© David Garlan 2009

38

Related work

- Many names for this area of research
 - self-healing systems, autonomic computing, self-* systems
 - Hot area: conferences, workshops, journals
- Most existing approaches focusing on specific domains and quality attributes
 - Usually assumes particular architectural style
- Other techniques used
 - Planning, expert systems
- Traditional areas
 - Control systems, utility theory, Markov decision processes, etc.

On-going work

- Cyber-physical systems
 - Large-scale distributed control systems
 - Examples: power grid, air-traffic control, multi-car coordination, smart home
- Analysis
 - How can you decide if a given adaptation strategy is "sound"?
 - Can model-checking be applied to Stitch?
- Service-oriented architectures
 - Adapting service orchestrations to match environment
- Incorporating learning
 - To improve strategy selection
 - To learn new strategies

Conclusion

- Today systems **fail to adapt** to changing environments
- **Architecture models** and an **adaptation language** can be combined to self-adapt systems for multiple quality dimensions with broad scope
- **Rainbow**
 - integrates architecture model and a language for self-adaptation
 - provides software engineers the ability to add and evolve self-adaptation capabilities
 - *cost-effectively*
 - *for multiple objectives*
 - *with explicit customization points to tailor self-adaptation capabilities for particular classes of systems and quality objectives.*

"Somewhere Over the Rainbow, A Stitch in Time Saves Nine"